From The Operator's Thesis: How Healthcare Really Buys AI (Saivai Group, 2026) — editable reader copy of the printed instrument; text identical to the book. © 2026 Anthony W. Russell. For the reader's own use; not legal, clinical, compliance, investment, accounting, tax, or consulting advice. Reader resource hub: saivaigroup.com/thesis

Instrument 7 — Judgment Allocation Map *(Chapter 7)*

For each task in the target workflow, assign a category and a supervision level. Category: machine-favored (high-volume, pattern-heavy, errors detectable and recoverable — run with audit); human-led, machine-assisted (machine drafts and narrows, human decides: the resident-and-attending structure); human-dominant (ambiguity, ethics, or accountability make delegation impossible: AI supports the surroundings only). Supervision level, per the Attending Model: observe; propose with co-signature; act with review; act with retrospective audit; act independently within a bounded domain: each granted by domain, earned by demonstrated performance, and revocable. Then name the two humans in the loop: the frontline lead who manages the process where it runs, and the accountable administrator who owns value, burden shift, compliance, and drift at the executive level. Finally, the training plan: no human enters the loop untrained — what does each role learn, when, and how is competence checked? If the vendor's answer to the human's role is "rescue" rather than "judgment," you are pricing a services business, not a safety architecture.